

Hepatitis C Virus Classification

Dataset: <https://archive.ics.uci.edu/dataset/571/hcv+data>

Libraries import

```
In [ ]: from utils import *
import matplotlib.pyplot as plt
import numpy as np
from sklearn.preprocessing import MinMaxScaler
from sklearn.calibration import cross_val_predict
from sklearn.neighbors import KNeighborsClassifier
from sklearn.neighbors import KNeighborsClassifier
from sklearn.ensemble import GradientBoostingClassifier
from sklearn.linear_model import LogisticRegression
from sklearn.calibration import cross_val_predict
import seaborn as sn
from sklearn.metrics import roc_curve, auc

def myPlotROCcurve(target_test, prediction, text=""):
    fpr, tpr, _ = roc_curve(target_test, prediction)
    roc_auc = auc(fpr, tpr)
    plt.figure(figsize=(8, 6))
    plt.plot(fpr, tpr, color='darkorange', lw=2, label='ROC curve (area = {:.2f})'.format(roc_auc))
    plt.plot([0, 1], [0, 1], color='navy', lw=2, linestyle='--')
    plt.xlabel('False Positive Rate')
    plt.ylabel('True Positive Rate')
    plt.title('Receiver Operating Characteristic (ROC) - ' + text)
    plt.legend(loc='lower right')
    plt.show()

def myPlotConfusionMatrix(target_test, prediction, text=""):
    conf_matrix = confusion_matrix(target_test, prediction)
    sns.heatmap(conf_matrix, annot=True, fmt='d', cmap='Blues', cbar=False)
    plt.xlabel('Predicted Class')
    plt.ylabel('True Class')
    plt.title(text)
    plt.show()

def myCrossValidation(classifier, cv_techique, features, target, text=""):
    accuracy = cross_val_score(
        classifier, # Classifier
        features, # Feature matrix
        target, # Target vector
        cv=cv_techique, # Cross-validation technique
        scoring="accuracy", # Loss function
        n_jobs=-1) # Use all CPU scores

    accuracy = np.mean(accuracy)
```

```

precision = cross_val_score(
    classifier, # Classifier
    features, # Feature matrix
    target, # Target vector
    cv=cv_technique, # Cross-validation technique
    scoring="precision", # Loss function
    n_jobs=-1) # Use all CPU scores

precision = np.mean(precision)

recall = cross_val_score(
    classifier, # Classifier
    features, # Feature matrix
    target, # Target vector
    cv=cv_technique, # Cross-validation technique
    scoring="recall", # Loss function
    n_jobs=-1) # Use all CPU scores

recall = np.mean(recall)

f1 = cross_val_score(
    classifier, # Classifier
    features, # Feature matrix
    target, # Target vector
    cv=cv_technique, # Cross-validation technique
    scoring="f1", # Loss function
    n_jobs=-1) # Use all CPU scores

f1 = np.mean(f1)

predictions = cross_val_predict(
    classifier,
    features,
    target,
    cv=cv_technique,
    method='predict')

return pd.DataFrame({"Accuracy": accuracy,
                     "Precision": precision,
                     "Recall": recall,
                     "F1": f1,
                     "Model": text},
                     index=[0]), predictions

def myResultFormalizer(report, text=""):

    return pd.DataFrame({"Accuracy": float(report['accuracy']),
                         "Precision": float(report['0']['precision']),
                         "Recall": float(report['0']['recall']),
                         "F1": float(report['0']['f1-score']),
                         "Model": text},
                         index=[0])

```

Data import and manipulation

```

In [ ]: dataframe = pd.read_csv('hcvdat0.csv')

dataframe.drop("Unnamed: 0", axis=1, inplace=True)

dataframe.dropna(inplace=True)

scale_mapper = {
    "0=Blood Donor": 0,
    "0s=suspect Blood Donor": 2,
    "1=Hepatitis": 1,
    "2=Fibrosis": 1,
    "3=Cirrhosis": 1,}

# Dividing blood types in two types
dataframe['Category'] = dataframe['Category'].replace(scale_mapper)

dataframe['Sex'] = dataframe['Sex'].replace({"m":0, "f":1})

# Removal of suspect blood donor because they are not useful for model
dataframe = dataframe[dataframe['Category'] != 2]

# Removal of wrong categorized data
dataframe = dataframe[dataframe['Category'].isin([0, 1])]
dataframe = dataframe[dataframe['Sex'].isin([0, 1])]

# Define the age ranges and labels for each category
age_bins = [0, 18, 30, 40, 50, 60, 70, 120] # Define the age bins
age_labels = [0, 1, 2, 3, 4, 5, 6] # Define the labels for each age group

# Categorize ages into age groups
dataframe['Age'] = pd.cut(dataframe['Age'], bins=age_bins, labels=age_labels)

dataframe.head(20)

dataframe.to_csv("wrangled_data.csv")

```

Outlier detection

```

In [ ]: # Create a box plot for each feature
plt.figure(figsize=(12, 6))
sns.boxplot(data=dataframe)
plt.title('Box plot of features')
plt.xticks(rotation=45)
plt.show()

fig, axs = plt.subplots(2, 2, figsize=(10,10), constrained_layout=True)
categorical = ['Age', 'Sex', 'Category']
for i, f in enumerate(categorical):
    sns.countplot(y=f, data=dataframe, ax=axs[i//2][i%2], order=dataframe[f].value_)

fig, axs = plt.subplots(3, 4, figsize=(15, 10), constrained_layout=True)
numerical = ['ALB', 'ALP', 'ALT', 'AST', 'BIL', 'CHE', 'CHOL', 'CREA', 'GGT', 'PROT']
for i, f in enumerate(numerical):

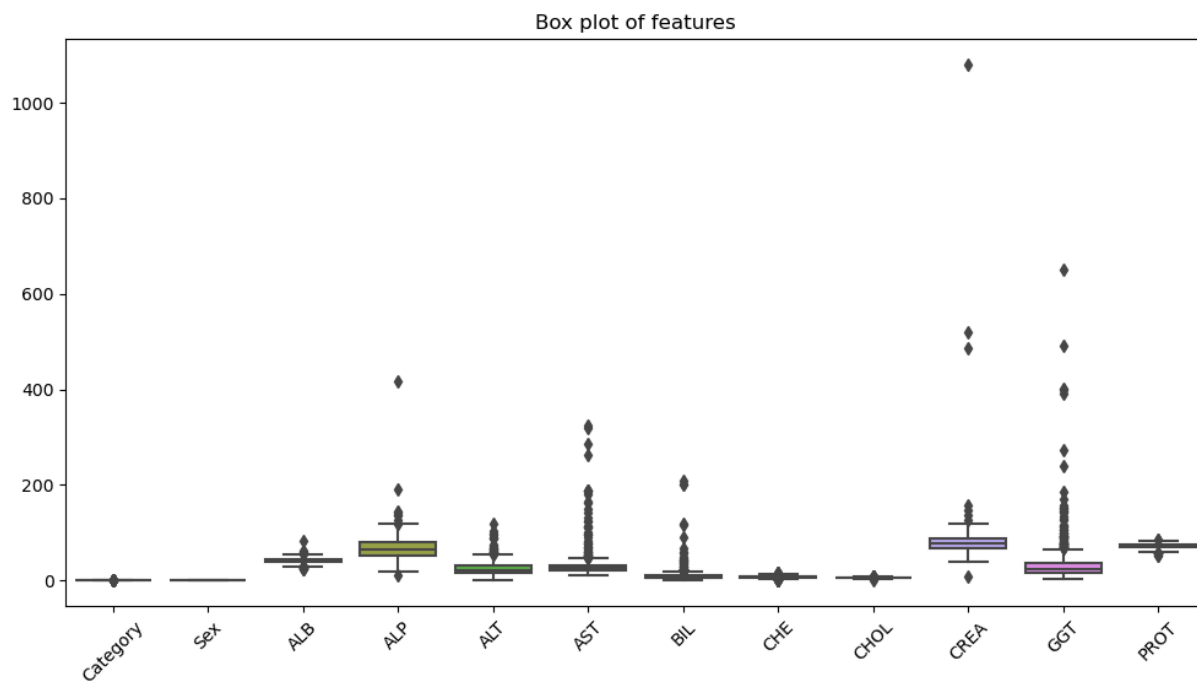
```

```





sns.histplot(x=f, data=dataframe, ax=axes[i//4][i%4], bins=100) # Adjusted index
plt.show()

corr_matrix = dataframe[ [*numerical, *categorical]].corr()
sns.heatmap(corr_matrix, annot=True)
plt.show()
print(corr_matrix)

```



```
c:\Users\Nemanja\miniconda3\envs\mlenv\lib\site-packages\seaborn\categorical.py:641:
FutureWarning: The default of observed=False is deprecated and will be changed to Tr
ue in a future version of pandas. Pass observed=False to retain current behavior or
observed=True to adopt the future default and silence this warning.
    grouped_vals = vals.groupby(grouper)
c:\Users\Nemanja\miniconda3\envs\mlenv\lib\site-packages\seaborn\_oldcore.py:1119: F
utureWarning: use_inf_as_na option is deprecated and will be removed in a future ver
sion. Convert inf values to NaN before operating instead.
    with pd.option_context('mode.use_inf_as_na', True):
c:\Users\Nemanja\miniconda3\envs\mlenv\lib\site-packages\seaborn\_oldcore.py:1119: F
utureWarning: use_inf_as_na option is deprecated and will be removed in a future ver
sion. Convert inf values to NaN before operating instead.
    with pd.option_context('mode.use_inf_as_na', True):
c:\Users\Nemanja\miniconda3\envs\mlenv\lib\site-packages\seaborn\_oldcore.py:1119: F
utureWarning: use_inf_as_na option is deprecated and will be removed in a future ver
sion. Convert inf values to NaN before operating instead.
    with pd.option_context('mode.use_inf_as_na', True):
c:\Users\Nemanja\miniconda3\envs\mlenv\lib\site-packages\seaborn\_oldcore.py:1119: F
utureWarning: use_inf_as_na option is deprecated and will be removed in a future ver
sion. Convert inf values to NaN before operating instead.
    with pd.option_context('mode.use_inf_as_na', True):
c:\Users\Nemanja\miniconda3\envs\mlenv\lib\site-packages\seaborn\_oldcore.py:1119: F
utureWarning: use_inf_as_na option is deprecated and will be removed in a future ver
sion. Convert inf values to NaN before operating instead.
    with pd.option_context('mode.use_inf_as_na', True):
c:\Users\Nemanja\miniconda3\envs\mlenv\lib\site-packages\seaborn\_oldcore.py:1119: F
utureWarning: use_inf_as_na option is deprecated and will be removed in a future ver
sion. Convert inf values to NaN before operating instead.
    with pd.option_context('mode.use_inf_as_na', True):
c:\Users\Nemanja\miniconda3\envs\mlenv\lib\site-packages\seaborn\_oldcore.py:1119: F
utureWarning: use_inf_as_na option is deprecated and will be removed in a future ver
sion. Convert inf values to NaN before operating instead.
    with pd.option_context('mode.use_inf_as_na', True):
c:\Users\Nemanja\miniconda3\envs\mlenv\lib\site-packages\seaborn\_oldcore.py:1119: F
utureWarning: use_inf_as_na option is deprecated and will be removed in a future ver
sion. Convert inf values to NaN before operating instead.
    with pd.option_context('mode.use_inf_as_na', True):
c:\Users\Nemanja\miniconda3\envs\mlenv\lib\site-packages\seaborn\_oldcore.py:1119: F
utureWarning: use_inf_as_na option is deprecated and will be removed in a future ver
sion. Convert inf values to NaN before operating instead.
    with pd.option_context('mode.use_inf_as_na', True):
```



	ALB	ALP	ALT	AST	BIL	CHE	\
ALB	1.000000	-0.103431	0.200235	-0.161389	-0.194822	0.364617	
ALP	-0.103431	1.000000	0.078155	0.027222	0.066974	0.043800	
ALT	0.200235	0.078155	1.000000	0.100191	-0.126918	0.311920	
AST	-0.161389	0.027222	0.100191	1.000000	0.321066	-0.224441	
BIL	-0.194822	0.066974	-0.126918	0.321066	1.000000	-0.330494	
CHE	0.364617	0.043800	0.311920	-0.224441	-0.330494	1.000000	
CHOL	0.167053	0.135502	0.184870	-0.211018	-0.187060	0.428312	
CREA	-0.018850	0.167393	-0.024710	-0.009114	0.019625	-0.013163	
GGT	-0.073379	0.428599	0.102898	0.481591	0.234191	-0.074410	
PROT	0.493732	-0.014861	0.146857	0.048419	-0.071996	0.297580	
Age	-0.159642	0.173311	-0.114250	0.057835	0.048684	-0.085646	
Sex	-0.178748	0.003135	-0.272588	-0.134649	-0.111291	-0.185438	
Category	-0.203710	-0.062378	-0.233696	0.645313	0.442584	-0.248236	

	CHOL	CREA	GGT	PROT	Age	Sex	Category
ALB	0.167053	-0.018850	-0.073379	0.493732	-0.159642	-0.178748	-0.203710
ALP	0.135502	0.167393	0.428599	-0.014861	0.173311	0.003135	-0.062378
ALT	0.184870	-0.024710	0.102898	0.146857	-0.114250	-0.272588	-0.233696
AST	-0.211018	-0.009114	0.481591	0.048419	0.057835	-0.134649	0.645313
BIL	-0.187060	0.019625	0.234191	-0.071996	0.048684	-0.111291	0.442584
CHE	0.428312	-0.013163	-0.074410	0.297580	-0.085646	-0.185438	-0.248236
CHOL	1.000000	-0.060087	0.031976	0.201525	0.137555	0.025227	-0.252205
CREA	-0.060087	1.000000	0.128460	-0.061710	-0.032940	-0.160133	0.166441
GGT	0.031976	0.128460	1.000000	0.049717	0.138731	-0.133407	0.461679
PROT	0.201525	-0.061710	0.049717	1.000000	-0.108158	-0.069882	0.006535
Age	0.137555	-0.032940	0.138731	-0.108158	1.000000	0.032236	0.037018
Sex	0.025227	-0.160133	-0.133407	-0.069882	0.032236	1.000000	-0.067596
Category	-0.252205	0.166441	0.461679	0.006535	0.037018	-0.067596	1.000000

```
In [ ]: columns_to_scale = ['ALB', 'ALP', 'ALT', 'AST', 'BIL', 'CHE', 'CHOL', 'CREA']

scaler = MinMaxScaler(feature_range=(-1, 1))
dataframe[columns_to_scale] = scaler.fit_transform(dataframe[columns_to_scale])

outlier_detector = EllipticEnvelope(contamination=.009)

# Fit detector
outlier_detector.fit(dataframe[columns_to_scale])

# Predict outliers
outliers = outlier_detector.predict(dataframe[columns_to_scale])
outliers_indices = outliers == -1
dataframe = dataframe[~outliers_indices]

# Create a box plot for each feature
plt.figure(figsize=(12, 6))
sns.boxplot(data=dataframe)
plt.title('Box plot of features')
plt.xticks(rotation=45)
plt.show()


fig, axs = plt.subplots(2, 2, figsize=(10,10), constrained_layout=True)
categorical = ['Age', 'Sex', 'Category']
for i, f in enumerate(categorical):
    sns.countplot(y=f, data=dataframe, ax=axs[i//2][i%2], order=dataframe[f].value_
```







```
fig, axs = plt.subplots(3, 4, figsize=(15, 10), constrained_layout=True)
numerical = ['ALB', 'ALP', 'ALB', 'ALT', 'AST', 'BIL', 'CHE', 'CHOL', 'CREA', 'GGT']
for i, f in enumerate(numerical):
    sns.histplot(x=f, data=dataframe, ax=axs[i//4][i%4], bins=100) # Adjusted index
plt.show()

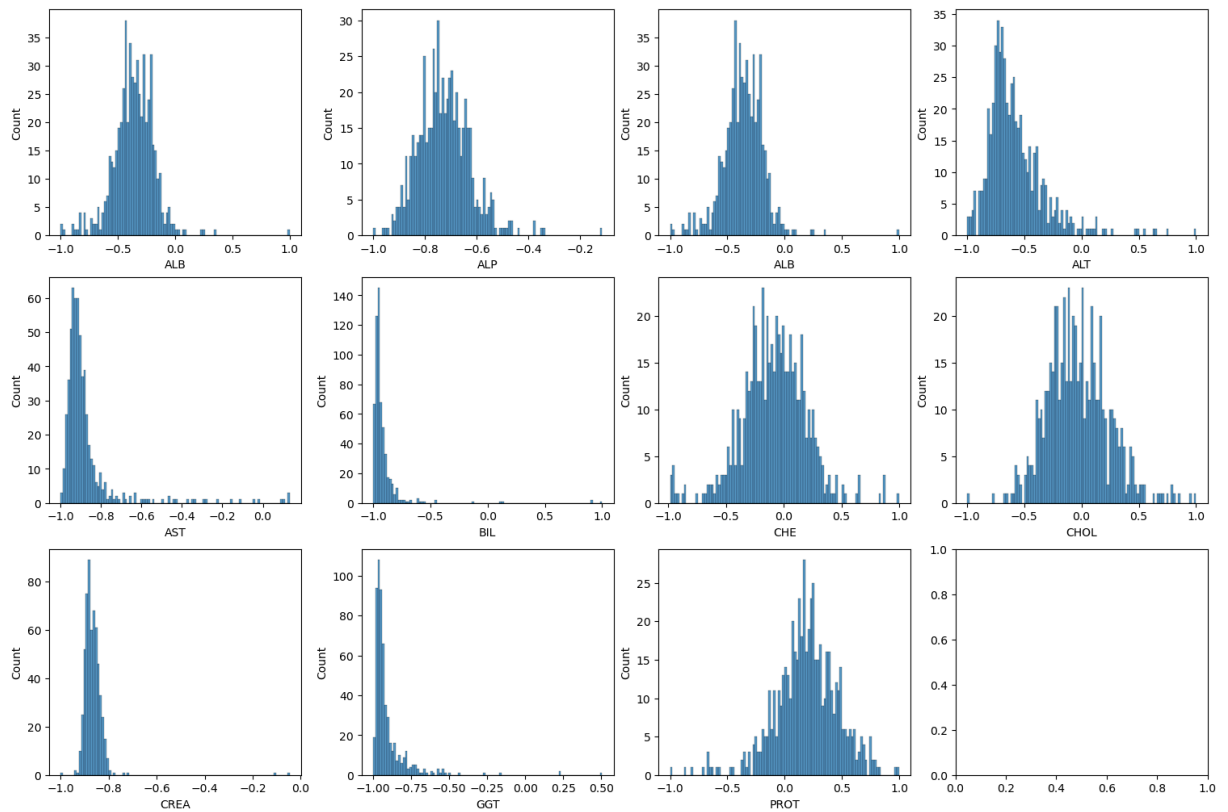
dataframe.head(15)
```

c:\Users\Nemanja\miniconda3\envs\mlenv\lib\site-packages\sklearn\covariance_robust_covariance.py:747: UserWarning: The covariance matrix associated to your dataset is not full rank
warnings.warn(



```
c:\Users\Nemanja\miniconda3\envs\mlenv\lib\site-packages\seaborn\categorical.py:641:
FutureWarning: The default of observed=False is deprecated and will be changed to Tr
ue in a future version of pandas. Pass observed=False to retain current behavior or
observed=True to adopt the future default and silence this warning.
    grouped_vals = vals.groupby(grouper)
c:\Users\Nemanja\miniconda3\envs\mlenv\lib\site-packages\seaborn\_oldcore.py:1119: F
utureWarning: use_inf_as_na option is deprecated and will be removed in a future ver
sion. Convert inf values to NaN before operating instead.
    with pd.option_context('mode.use_inf_as_na', True):
c:\Users\Nemanja\miniconda3\envs\mlenv\lib\site-packages\seaborn\_oldcore.py:1119: F
utureWarning: use_inf_as_na option is deprecated and will be removed in a future ver
sion. Convert inf values to NaN before operating instead.
    with pd.option_context('mode.use_inf_as_na', True):
c:\Users\Nemanja\miniconda3\envs\mlenv\lib\site-packages\seaborn\_oldcore.py:1119: F
utureWarning: use_inf_as_na option is deprecated and will be removed in a future ver
sion. Convert inf values to NaN before operating instead.
    with pd.option_context('mode.use_inf_as_na', True):
c:\Users\Nemanja\miniconda3\envs\mlenv\lib\site-packages\seaborn\_oldcore.py:1119: F
utureWarning: use_inf_as_na option is deprecated and will be removed in a future ver
sion. Convert inf values to NaN before operating instead.
    with pd.option_context('mode.use_inf_as_na', True):
c:\Users\Nemanja\miniconda3\envs\mlenv\lib\site-packages\seaborn\_oldcore.py:1119: F
utureWarning: use_inf_as_na option is deprecated and will be removed in a future ver
sion. Convert inf values to NaN before operating instead.
    with pd.option_context('mode.use_inf_as_na', True):
c:\Users\Nemanja\miniconda3\envs\mlenv\lib\site-packages\seaborn\_oldcore.py:1119: F
utureWarning: use_inf_as_na option is deprecated and will be removed in a future ver
sion. Convert inf values to NaN before operating instead.
    with pd.option_context('mode.use_inf_as_na', True):
c:\Users\Nemanja\miniconda3\envs\mlenv\lib\site-packages\seaborn\_oldcore.py:1119: F
utureWarning: use_inf_as_na option is deprecated and will be removed in a future ver
sion. Convert inf values to NaN before operating instead.
    with pd.option_context('mode.use_inf_as_na', True):
c:\Users\Nemanja\miniconda3\envs\mlenv\lib\site-packages\seaborn\_oldcore.py:1119: F
utureWarning: use_inf_as_na option is deprecated and will be removed in a future ver
sion. Convert inf values to NaN before operating instead.
    with pd.option_context('mode.use_inf_as_na', True):
c:\Users\Nemanja\miniconda3\envs\mlenv\lib\site-packages\seaborn\_oldcore.py:1119: F
utureWarning: use_inf_as_na option is deprecated and will be removed in a future ver
sion. Convert inf values to NaN before operating instead.
    with pd.option_context('mode.use_inf_as_na', True):
c:\Users\Nemanja\miniconda3\envs\mlenv\lib\site-packages\seaborn\_oldcore.py:1119: F
utureWarning: use_inf_as_na option is deprecated and will be removed in a future ver
sion. Convert inf values to NaN before operating instead.
    with pd.option_context('mode.use_inf_as_na', True):
```





Out[]:

	Category	Age	Sex	ALB	ALP	ALT	AST	BIL	CHE	
0	0	2	0	-0.476351	-0.796694	-0.883959	-0.935256	-0.944981	-0.264843	-
1	0	2	0	-0.476351	-0.708858	-0.708191	-0.918590	-0.979730	0.300867	-
2	0	2	0	-0.192568	-0.687145	-0.397611	-0.739744	-0.958494	-0.010007	-
3	0	2	0	-0.317568	-0.799161	-0.493174	-0.932051	-0.834942	-0.211474	-
4	0	2	0	-0.452703	-0.690106	-0.459044	-0.917949	-0.924710	0.031354	-
5	0	2	0	-0.371622	-0.842092	-0.699659	-0.950641	-0.898649	0.134089	-
6	0	2	0	-0.212838	-0.851962	-0.716724	-0.962821	-0.935328	-0.254169	-
7	0	2	0	-0.351351	-0.849001	-0.404437	-0.877564	-0.861969	-0.412942	-
8	0	2	0	-0.057432	-0.732544	-0.619454	-0.941026	-0.950772	-0.030020	-
9	0	2	0	-0.344595	-0.629904	-0.668942	-0.948718	-0.677606	-0.460974	-
10	0	2	0	-0.280405	-0.797681	-0.645051	-0.933333	-0.851351	-0.635757	-
11	0	2	0	-0.209459	-0.719220	-0.839590	-0.948718	-0.962355	-0.207472	-
12	0	2	0	-0.550676	-0.667900	-0.612628	-0.935897	-0.949807	-0.047365	-
13	0	2	0	-0.459459	-0.800642	-0.744027	-0.923077	-0.951737	-0.327552	-
14	0	2	0	-0.469595	-0.859363	-0.631399	-0.929487	-0.977799	-0.571714	-

Train-Test splitting

```
In [ ]: dataframe_shuffled = dataframe.sample(frac=1)

features, target = dataframe_shuffled.drop('Category', axis=1), dataframe_shuffled['Category']

# Split into training and test set
features_train, features_test, target_train, target_test = train_test_split(
    features, target, random_state=0)

# List that contains results of each model
results_list = list()
results_list_cv_only = list()
```

Model testing

Dummy classifier

```
In [ ]: # Create dummy classifier
dummy = DummyClassifier(strategy='uniform', random_state=1)

# "Train" model
dummy.fit(features_train, target_train)

# Predict on test features
dummy_prediction = dummy.predict(features_test)

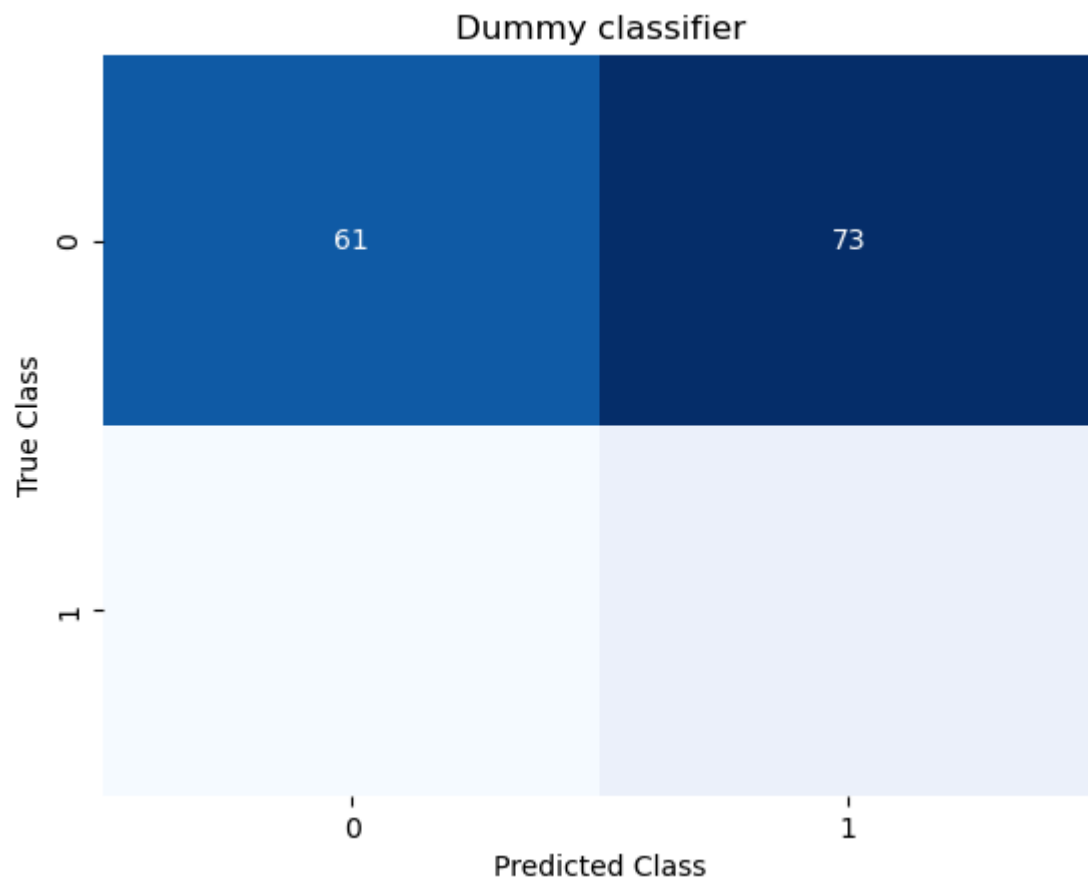
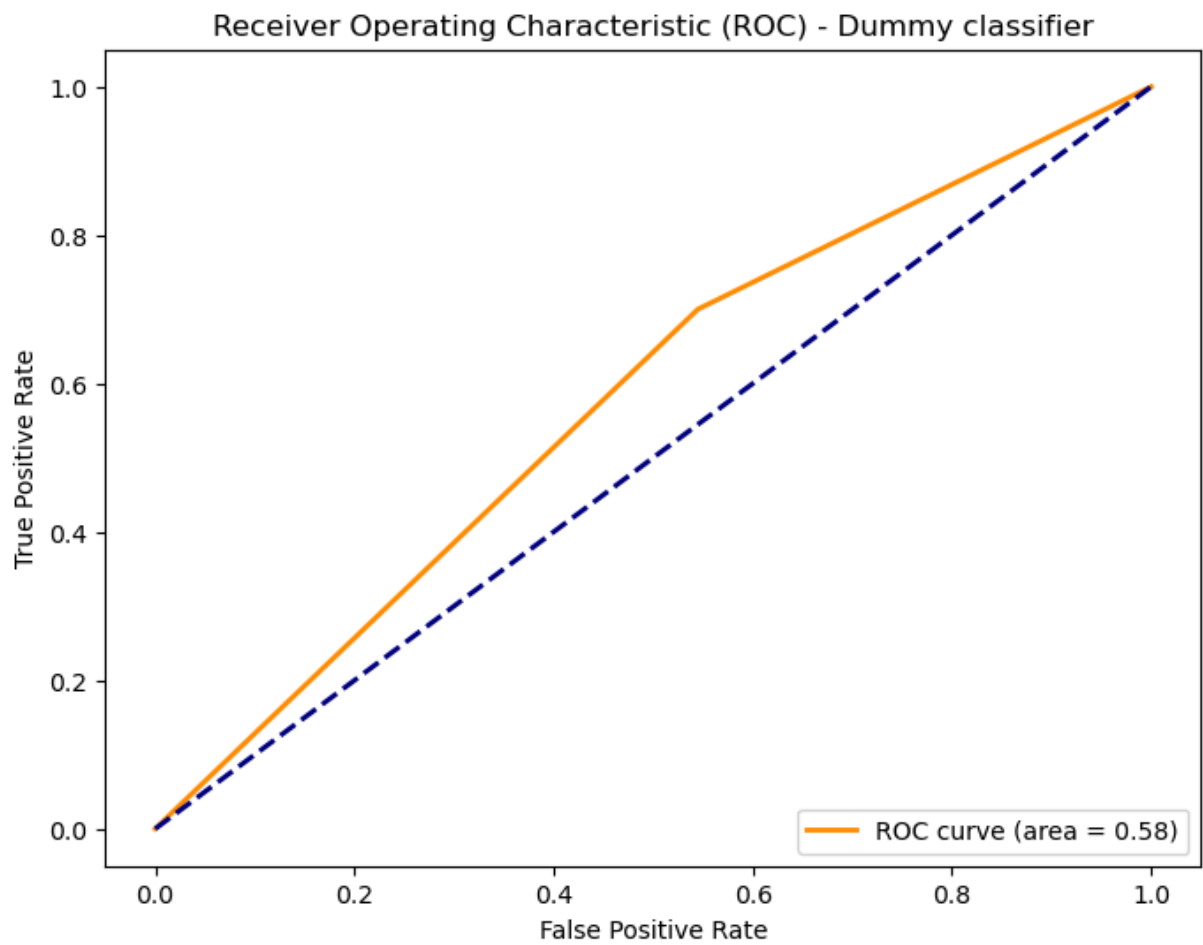
# Visualize the ROC curve
myPlotROCcurve(target_test, dummy_prediction, "Dummy classifier")

# Visualize the confusion matrix
myPlotConfusionMatrix(target_test, dummy_prediction, "Dummy classifier")

dummy_report = classification_report(target_test, dummy_prediction, output_dict=True)
print(classification_report(target_test, dummy_prediction))

dummy_bs_report_fromalized = myResultFormalizer(dummy_report, "Dummy classifier")

results_list.append(dummy_bs_report_fromalized)
```



	precision	recall	f1-score	support
0	0.95	0.46	0.62	134
1	0.09	0.70	0.16	10
accuracy			0.47	144
macro avg	0.52	0.58	0.39	144
weighted avg	0.89	0.47	0.58	144

RandomForestClassifier, Basic split

```
In [ ]: # Create classifier
rf_classifier = RandomForestClassifier()

model_title = "Random Forest Classifier, Basic split"

# Train model
rf_classifier.fit(features_train, target_train)

# Predict on test features
rfc_prediction = rf_classifier.predict(features_test)



# Visualize the ROC curve
myPlotROCcurve(target_test, rfc_prediction)

# Visualize the confusion matrix
myPlotConfusionMatrix(target_test, rfc_prediction)

# Print classification report
print("\nBasic split Random Forest Classifier Classification Report:")
rfc_report = classification_report(target_test, rfc_prediction, output_dict=True)
print(classification_report(target_test, rfc_prediction))

rfc_bs_report_formalized = myResultFormalizer(rfc_report, model_title)

results_list.append(rfc_bs_report_formalized)
```



Basic split Random Forest Classifier Classification Report:

	precision	recall	f1-score	support
0	1.00	1.00	1.00	134
1	1.00	1.00	1.00	10
accuracy			1.00	144
macro avg	1.00	1.00	1.00	144
weighted avg	1.00	1.00	1.00	144

RandomForestClassifier, Cross-Validation

```
In [ ]: # Create Random Forest Classifier object
rf_classifier = RandomForestClassifier()

model_title = "Random Forest Classifier and Cross validation"

# Stratified K-Fold cross-validation
skf = StratifiedKFold(n_splits=10, shuffle=True, random_state=1)



rf_cv_results, predictions = myCrossValidation(rf_classifier, skf, features, target)

myPlotROCCurve(target, predictions)

myPlotConfusionMatrix(target, predictions, model_title)

rf_cv_results.head()

results_list.append(rf_cv_results)
results_list_cv_only.append(rf_cv_results)
```



Naive Bayes, Basic split

```
In [ ]: # Create Naive Bayes Classifier object
nb_classifier = GaussianNB()

model_title = "Naive Bayes, Basic split"

# Train model
nb_classifier.fit(features_train, target_train)

# Predict on test features
nb_prediction = nb_classifier.predict(features_test)

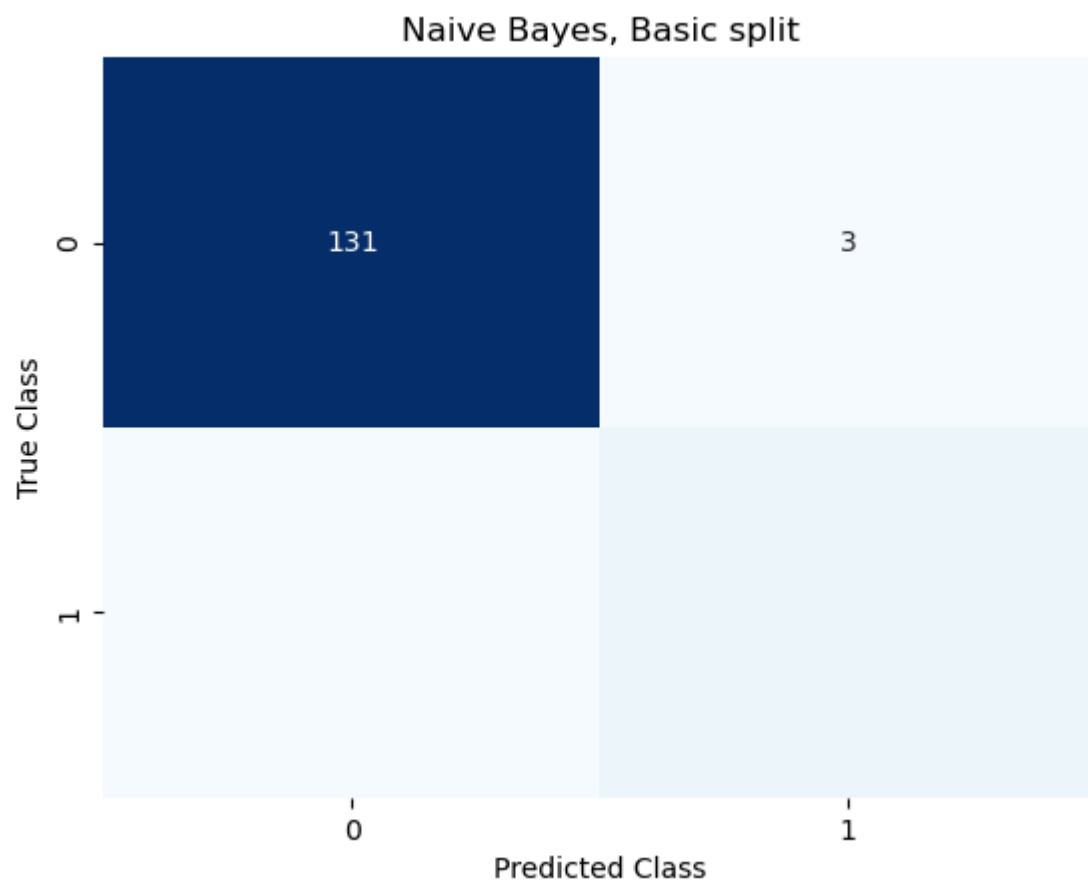

# Visualize the ROC curve
myPlotROCcurve(target_test, nb_prediction, model_title)

# Visualize the confusion matrix
myPlotConfusionMatrix(target_test, nb_prediction, model_title)

# Print classification report
print("\nBasic split Naive Bayes Classifier Classification Report:")
nb_report = classification_report(target_test, nb_prediction, output_dict=True)
print(classification_report(target_test, nb_prediction))

nb_bs_report_formalized = myResultFormalizer(nb_report, model_title)

results_list.append(nb_bs_report_formalized)
```



Basic split Naive Bayes Classifier Classification Report:

	precision	recall	f1-score	support
0	0.98	0.98	0.98	134
1	0.73	0.80	0.76	10
accuracy			0.97	144
macro avg	0.86	0.89	0.87	144
weighted avg	0.97	0.97	0.97	144

Naive Bayes, Cross-Validation

```
In [ ]: # Create Naive Bayes Classifier object
nb_classifier = GaussianNB()

model_title = "Naive Bayes Classifier and Cross validation"

# Stratified K-Fold cross-validation
skf = StratifiedKFold(n_splits=10, shuffle=True, random_state=1)

nb_cv_results, predictions = myCrossValidation(nb_classifier, skf, features, target)


myPlotROCCurve(target, predictions, model_title)

myPlotConfusionMatrix(target, predictions, model_title)


nb_cv_results.head()

results_list.append(nb_cv_results)
results_list_cv_only.append(nb_cv_results)
```

Receiver Operating Characteristic (ROC) - Naive Bayes Classifier and Cross validation



Naive Bayes Classifier and Cross validation



K-Nearest Neighbors, Basic split

```

In [ ]: # Create K-Nearest Neighbors classifier object
knn_classifier = KNeighborsClassifier()

model_title = "K-Nearest Neighbours, Basic split"

# Train model
knn_classifier.fit(features_train, target_train)

# Predict on test features
knn_prediction = knn_classifier.predict(features_test)

# Visualize the ROC curve
myPlotROCcurve(target_test, knn_prediction, model_title)

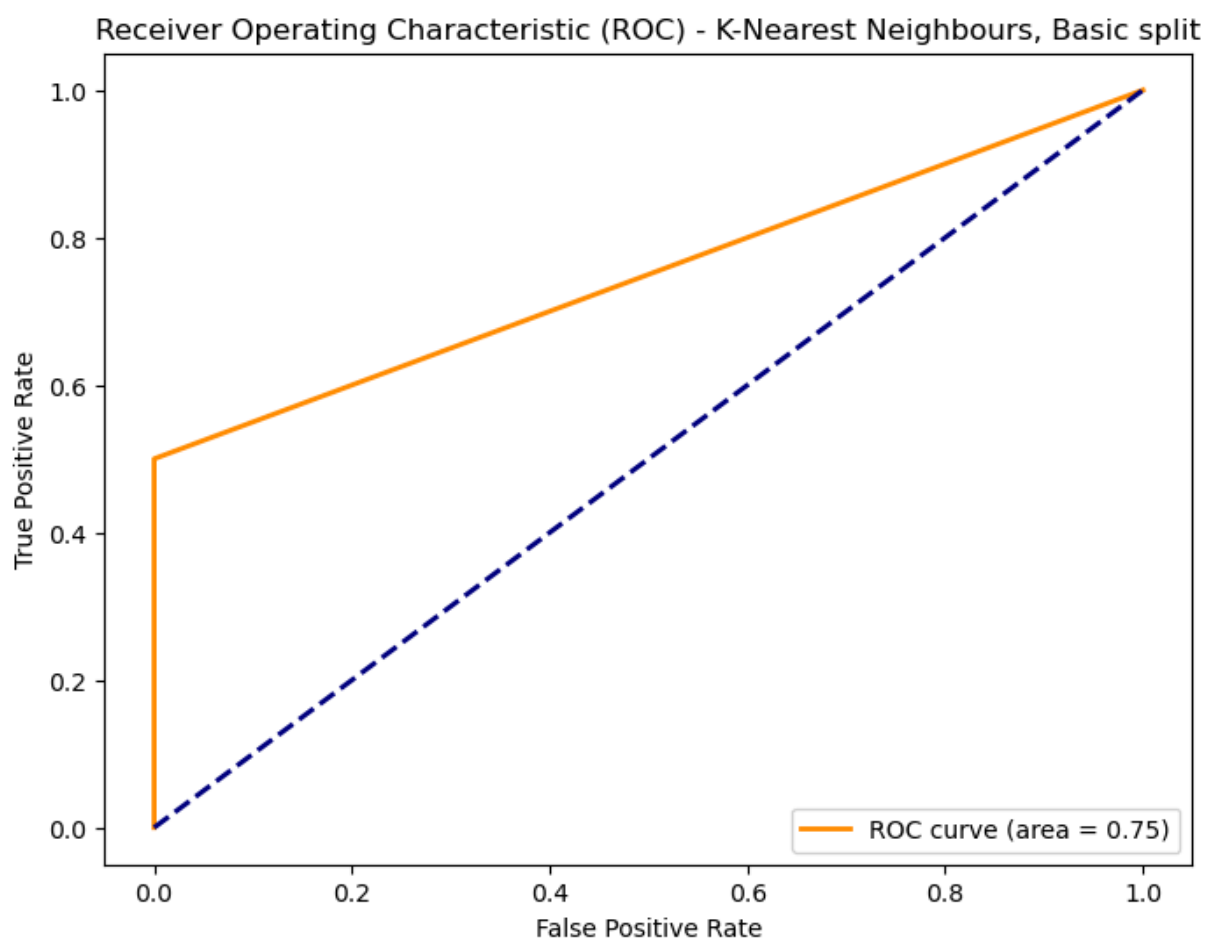
# Visualize the confusion matrix
myPlotConfusionMatrix(target_test, knn_prediction, model_title)

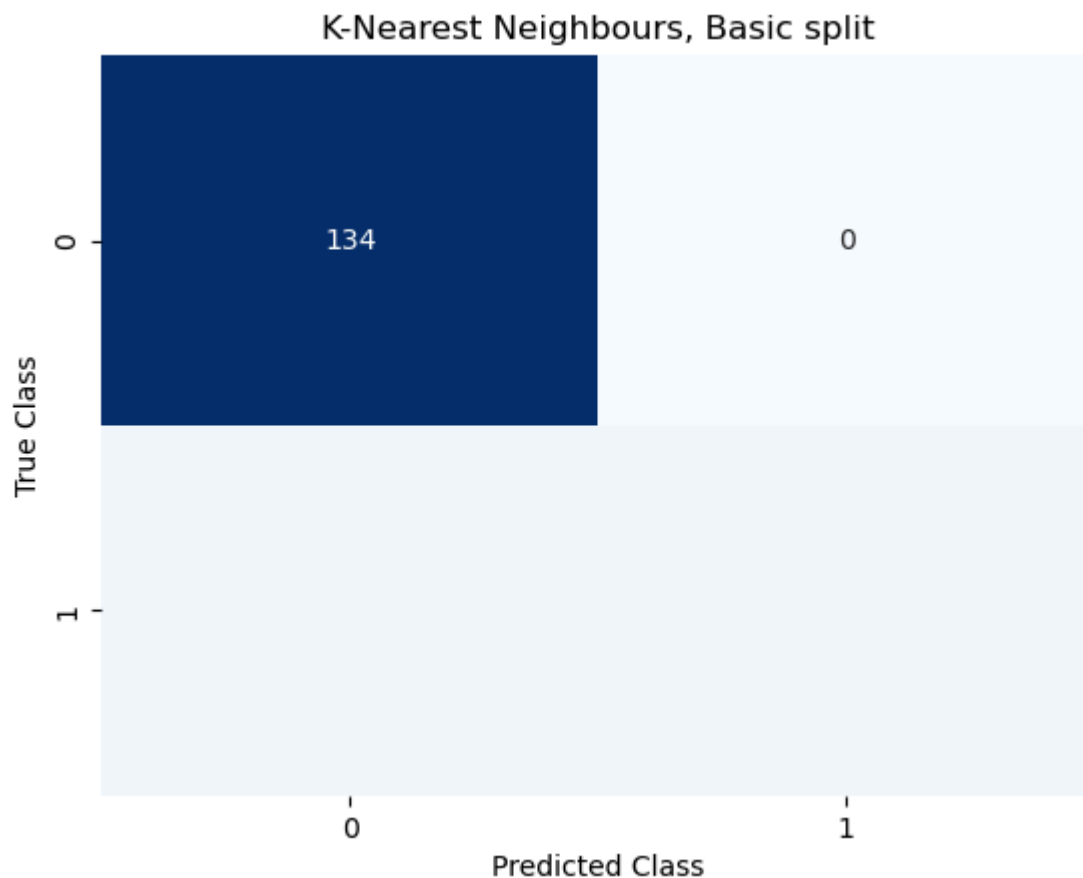
# Print classification report
print("\nBasic split K-Nearest Neighbours Classifier Classification Report:")
knn_bs_report = classification_report(target_test, knn_prediction, output_dict=True)
print(classification_report(target_test, knn_prediction))

knn_bs_report_formalized = myResultFormalizer(knn_bs_report, model_title)

results_list.append(knn_bs_report_formalized)

```





Basic split K-Nearest Neighbours Classifier Classification Report:

	precision	recall	f1-score	support
0	0.96	1.00	0.98	134
1	1.00	0.50	0.67	10
accuracy			0.97	144
macro avg	0.98	0.75	0.82	144
weighted avg	0.97	0.97	0.96	144

K-Nearest Neighbors, Cross-Validation

```
In [ ]: # Create K-Nearest Neighbors Classifier object
knn_classifier = KNeighborsClassifier()

model_title = "K-Nearest Neighbors Classifier and Cross validation"

# Stratified K-Fold cross-validation
skf = StratifiedKFold(n_splits=10, shuffle=True, random_state=1)

knn_cv_results, predictions = myCrossValidation(knn_classifier, skf, features, target)

myPlotROCCurve(target, predictions, model_title)


myPlotConfusionMatrix(target, predictions, model_title)

knn_cv_results.head()
```




```
results_list.append(knn_cv_results)
results_list_cv_only.append(knn_cv_results)
```

Receiver Operating Characteristic (ROC) - K-Nearest Neighbors Classifier and Cross validation



K-Nearest Neighbors Classifier and Cross validation



Gradient Boosting, Basic split

```
In [ ]: # Create Gradient Boosting classifier object
gb_classifier = GradientBoostingClassifier()

model_title = "Gradient Boosting, Basic split"

# Train model
gb_classifier.fit(features_train, target_train)

# Predict on test features
gb_prediction = gb_classifier.predict(features_test)

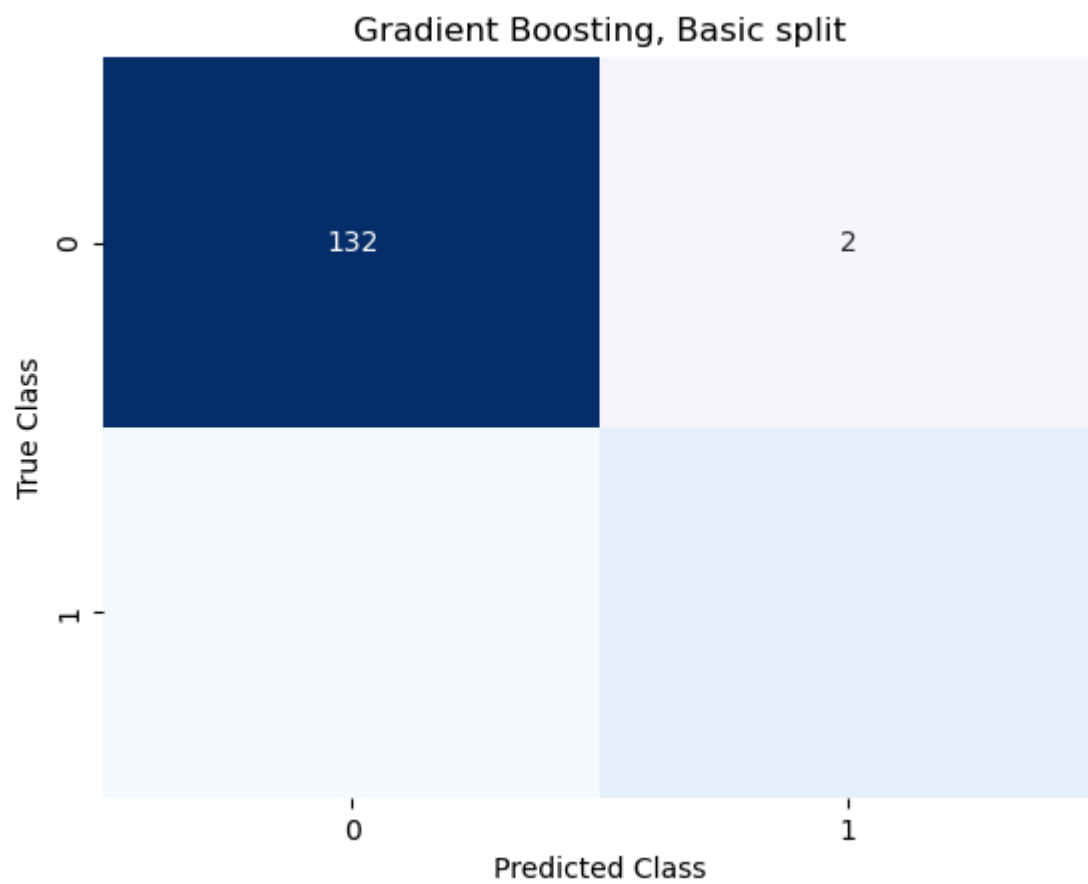
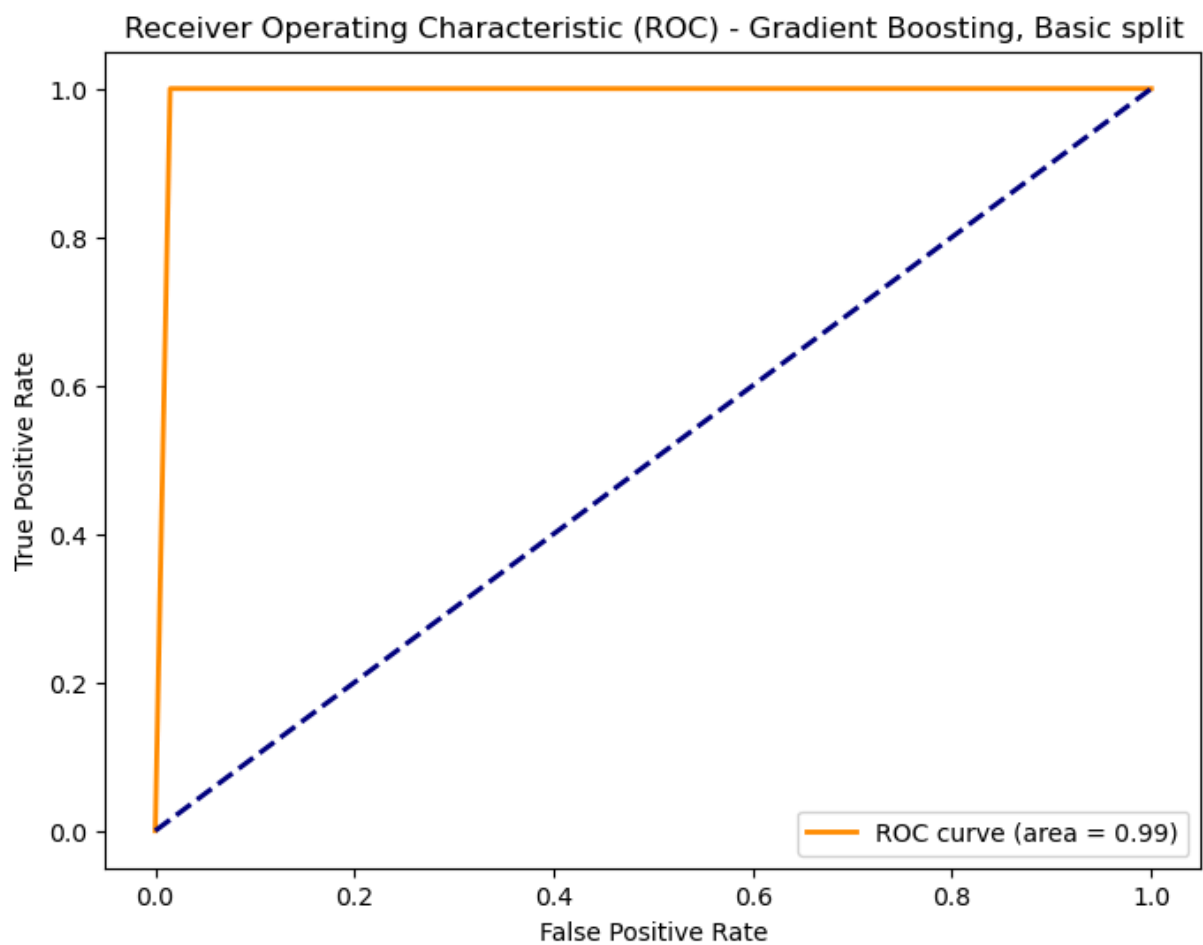
# Visualize the ROC curve
myPlotROCcurve(target_test, gb_prediction, model_title)

# Visualize the confusion matrix
myPlotConfusionMatrix(target_test, gb_prediction, model_title)

# Print classification report
print("\nBasic split Gradient Boosting Classifier Classification Report:")
gb_bs_report = classification_report(target_test, gb_prediction, output_dict=True)
print(classification_report(target_test, gb_prediction))

gb_bs_report_formalized = myResultFormalizer(gb_bs_report, model_title)

results_list.append(gb_bs_report_formalized)
```



Basic split Gradient Boosting Classifier Classification Report:

	precision	recall	f1-score	support
0	1.00	0.99	0.99	134
1	0.83	1.00	0.91	10
accuracy			0.99	144
macro avg	0.92	0.99	0.95	144
weighted avg	0.99	0.99	0.99	144

Gradient Boosting, Cross-Validation

```
In [ ]: # Create Gradient Boosting Classifier object
gb_classifier = GradientBoostingClassifier()

model_title = "Gradient Boosting Classifier and Cross validation"

# Stratified K-Fold cross-validation
skf = StratifiedKFold(n_splits=10, shuffle=True, random_state=1)

gb_cv_results, predictions = myCrossValidation(gb_classifier, skf, features, target)


myPlotROCCurve(target, predictions, model_title)

myPlotConfusionMatrix(target, predictions, model_title)

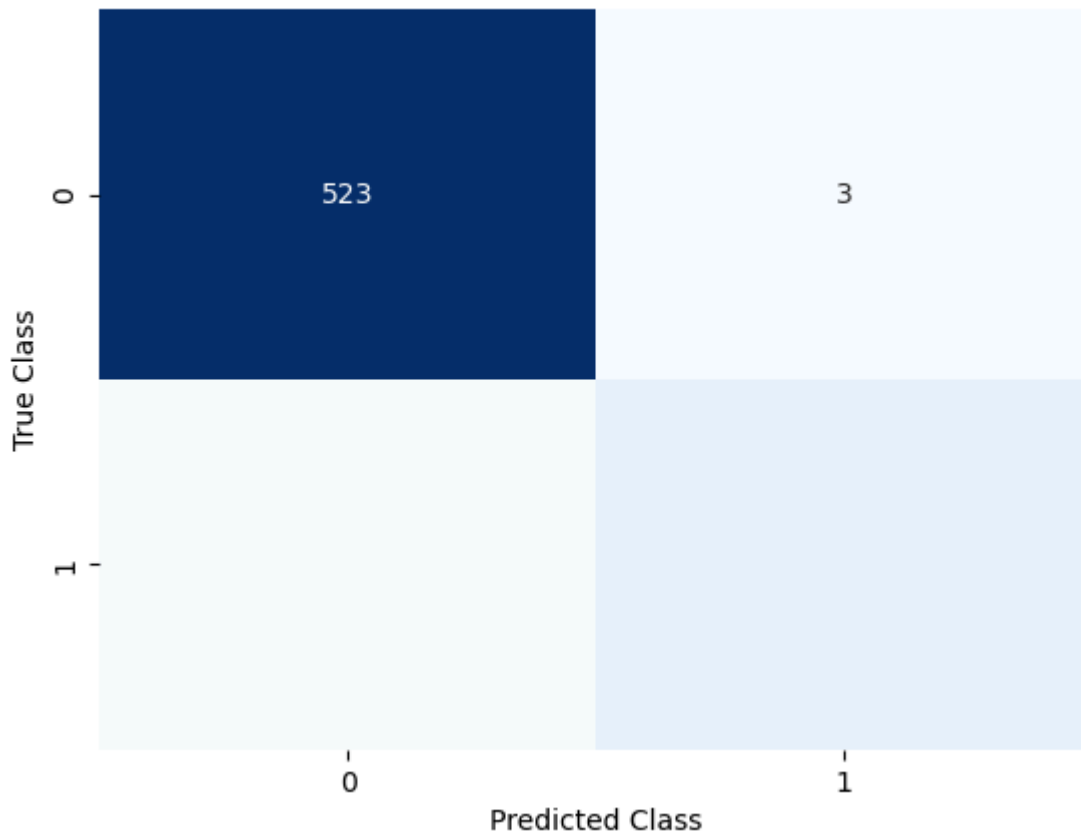
gb_cv_results.head()

results_list.append(gb_cv_results)
results_list_cv_only.append(gb_cv_results)
```

Receiver Operating Characteristic (ROC) - Gradient Boosting Classifier and Cross validation



Gradient Boosting Classifier and Cross validation



Logistic Regression, Basic split

```

In [ ]: # Create Logistic Regression classifier object
lr_classifier = LogisticRegression()

model_title = "Logistic Regression, Basic split"

# Train model
lr_classifier.fit(features_train, target_train)

# Predict on test features
lr_prediction = lr_classifier.predict(features_test)

# Visualize the ROC curve
myPlotROCcurve(target_test, lr_prediction, model_title)

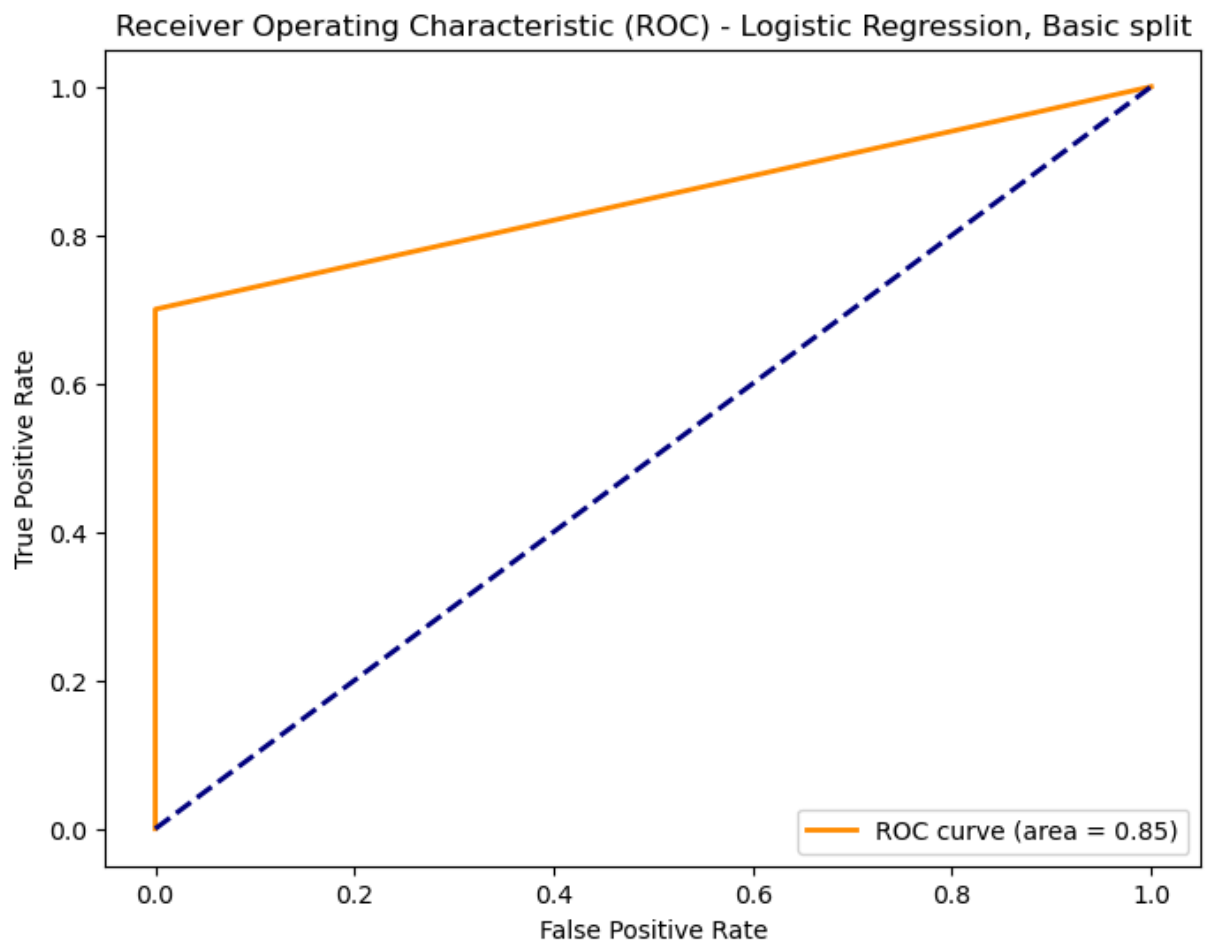
# Visualize the confusion matrix
myPlotConfusionMatrix(target_test, lr_prediction, model_title)

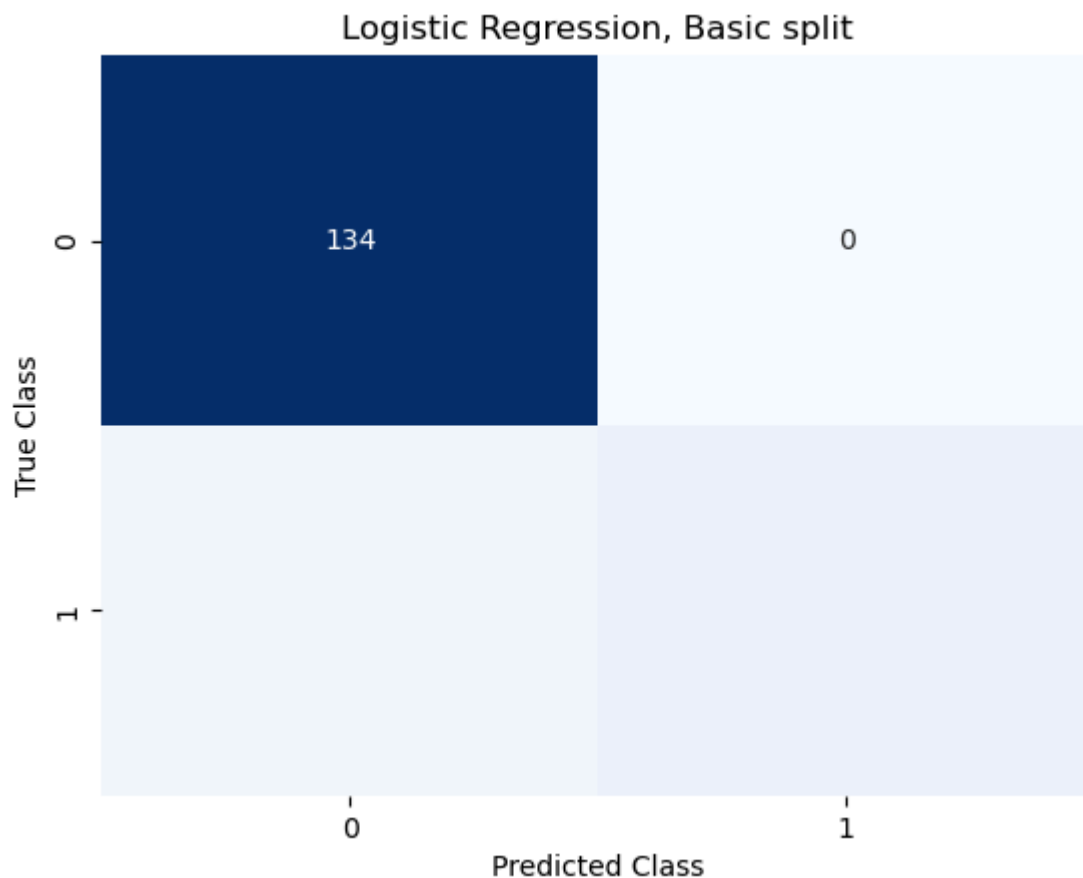
# Print classification report
print("\nBasic split Logistic Regression Classifier Classification Report:")
lr_bs_report = classification_report(target_test, lr_prediction, output_dict=True)
print(classification_report(target_test, lr_prediction))

gb_bs_report_formalized = myResultFormalizer(lr_bs_report, model_title)

results_list.append(gb_bs_report_formalized)

```





Basic split Logistic Regression Classifier Classification Report:

	precision	recall	f1-score	support
0	0.98	1.00	0.99	134
1	1.00	0.70	0.82	10
accuracy			0.98	144
macro avg	0.99	0.85	0.91	144
weighted avg	0.98	0.98	0.98	144

Logistic Regression, Cross-Validation

```
In [ ]: # Create Logistic Regression classifier object
lr_classifier_cv = LogisticRegression()

model_title = "Logistic Regression Classifier and Cross validation"

lr_cv_results, predictions = myCrossValidation(lr_classifier_cv, skf, features, tar


myPlotROCcurve(target, predictions, model_title)

myPlotConfusionMatrix(target, predictions, model_title)


lr_cv_results.head()

results_list.append(lr_cv_results)
results_list_cv_only.append(lr_cv_results)
```

Receiver Operating Characteristic (ROC) - Logistic Regression Classifier and Cross validation



Logistic Regression Classifier and Cross validation



Results compared


```
In [ ]: def plot_classifier_scores(dataframes, metrics, angle = 90):
    # Plotting
    bar_width = 0.15
    index = np.arange(len(dataframes))

    fig, ax = plt.subplots(figsize=(15, 8))

    for i, metric in enumerate(metrics):
        scores = [df[metric].iloc[0] for df in dataframes]
        bars = ax.bar(index + (i - len(metrics) / 2) * bar_width, scores, bar_width)

        # Add value labels on top of each bar
        for bar in bars:
            height = bar.get_height()
            ax.text(bar.get_x() + bar.get_width()/2, height, round(height, 2), ha='center')


    ax.set_xlabel('Classifiers')
    ax.set_ylabel('Scores')
    ax.set_title('Comparison of Scores')
    ax.set_xticks(index)
    ax.set_xticklabels([df['Model'].iloc[0] for df in dataframes], rotation=angle)
    ax.legend()


    plt.tight_layout()
    plt.show()

# List of feature names
feature_names = ["Accuracy", "Precision", "Recall", "F1"]

plot_classifier_scores(results_list, feature_names)

plot_classifier_scores(results_list_cv_only, feature_names)
```





Hiperparameters tuning

Nećemo razmatrati rezultate dobijene klasičnom podelom na train i test skupove zato što je dataset jako nebalansiran i rezultati koji su dobijeni nisu od relevantne koristi.

Za dalje podešavanje hiperparametara nastavljamo sa Random Forest i Gradient Boosting klasifikatorima zato što su dali najbolje rezultate.

Random Forest Hiperparameters tuning

```
In [ ]: from sklearn.ensemble import RandomForestClassifier
from sklearn.model_selection import GridSearchCV, RandomizedSearchCV

# Create a Random Forest classifier object
rf_classifier = RandomForestClassifier()

# Define the hyperparameter grid
param_grid = {
    'n_estimators': [10, 50, 100, 200],
    'max_depth': [None, 5, 10, 20],
    'min_samples_split': [2, 5, 10],
    'min_samples_leaf': [1, 2, 4],
    'max_features': ['auto', 'sqrt', 'log2', None],
    'bootstrap': [True, False]
}

# Perform randomized search
np.random.seed(42)
random_search = RandomizedSearchCV(estimator=rf_classifier, param_distributions=param_grid, n_iter=10)
random_search.fit(features, target)

rand_params = random_search.best_params_
```

```

## Perform grid search
#grid_search = GridSearchCV(estimator=rf_classifier, param_grid=param_grid, cv=5, n
#grid_search.fit(features, target)
## Best hyperparameters: {'bootstrap': True, 'max_depth': 20, 'max_features': 'log2

## Get the best hyperparameters
# best_params = grid_search.best_params_
# print("Best hyperparameters:", best_params)

best_params = {'bootstrap': True,
               'max_depth': 20,
               'max_features': 'log2',
               'min_samples_leaf': 1,
               'min_samples_split': 5,
               'n_estimators': 50}

best_params_classifier = RandomForestClassifier(**best_params)

model_title = "Random Forest Classifier and Cross validation, Hyperparameters tuned
rf_cv_hp_grid_results, predictions = myCrossValidation(best_params_classifier, skf, f
myPlotROCCurve(target, predictions, model_title)

myPlotConfusionMatrix(target, predictions, model_title)

rf_cv_hp_grid_results.head()

rnd_params_classifier = RandomForestClassifier(**rand_params)

model_title = "Random Forest Classifier and Cross validation, Hyperparameters tuned
rf_cv_hp_rnd_results, predictions = myCrossValidation(rnd_params_classifier, skf, f
myPlotROCCurve(target, predictions, model_title)

myPlotConfusionMatrix(target, predictions, model_title)

rf_cv_hp_rnd_results.head()

final_results = list()
final_results.append(rf_cv_hp_grid_results)
final_results.append(rf_cv_hp_rnd_results)

```

```
c:\Users\Nemanja\miniconda3\envs\mlenv\lib\site-packages\sklearn\model_selection\_validation.py:425: FitFailedWarning:
5 fits failed out of a total of 50.
The score on these train-test partitions for these parameters will be set to nan.
If these failures are not expected, you can try to debug them by setting error_score='raise'.
```

Below are more details about the failures:

4 fits failed with the following error:

Traceback (most recent call last):

```
File "c:\Users\Nemanja\miniconda3\envs\mlenv\lib\site-packages\sklearn\model_selection\_validation.py", line 732, in _fit_and_score
    estimator.fit(X_train, y_train, **fit_params)
File "c:\Users\Nemanja\miniconda3\envs\mlenv\lib\site-packages\sklearn\base.py", line 1144, in wrapper
    estimator._validate_params()
File "c:\Users\Nemanja\miniconda3\envs\mlenv\lib\site-packages\sklearn\base.py", line 637, in _validate_params
    validate_parameter_constraints(
File "c:\Users\Nemanja\miniconda3\envs\mlenv\lib\site-packages\sklearn\utils\_param_validation.py", line 95, in validate_parameter_constraints
    raise InvalidParameterError(
sklearn.utils._param_validation.InvalidParameterError: The 'max_features' parameter of RandomForestClassifier must be an int in the range [1, inf), a float in the range (0.0, 1.0], a str among {'sqrt', 'log2'} or None. Got 'auto' instead.
```


1 fits failed with the following error:

Traceback (most recent call last):


```
File "c:\Users\Nemanja\miniconda3\envs\mlenv\lib\site-packages\sklearn\model_selection\_validation.py", line 732, in _fit_and_score
    estimator.fit(X_train, y_train, **fit_params)
File "c:\Users\Nemanja\miniconda3\envs\mlenv\lib\site-packages\sklearn\base.py", line 1144, in wrapper
    estimator._validate_params()
File "c:\Users\Nemanja\miniconda3\envs\mlenv\lib\site-packages\sklearn\base.py", line 637, in _validate_params
    validate_parameter_constraints(
File "c:\Users\Nemanja\miniconda3\envs\mlenv\lib\site-packages\sklearn\utils\_param_validation.py", line 95, in validate_parameter_constraints
    raise InvalidParameterError(
sklearn.utils._param_validation.InvalidParameterError: The 'max_features' parameter of RandomForestClassifier must be an int in the range [1, inf), a float in the range (0.0, 1.0], a str among {'log2', 'sqrt'} or None. Got 'auto' instead.
```

```
warnings.warn(some_fits_failed_message, FitFailedWarning)
c:\Users\Nemanja\miniconda3\envs\mlenv\lib\site-packages\sklearn\model_selection\_search.py:976: UserWarning: One or more of the test scores are non-finite: [0.97742129 0.97742129 0.97916042 0.98088456 0.98262369 0.98089955
nan 0.98088456 0.98088456 0.97742129]
warnings.warn(
```


Receiver Operating Characteristic (ROC) - Random Forest Classifier and Cross validation, Hyperparameters tuned (Grid)




Random Forest Classifier and Cross validation, Hyperparameters tuned (Grid)



Receiver Operating Characteristic (ROC) - Random Forest Classifier and Cross validation, Hyperparameters tuned (Random)



Random Forest Classifier and Cross validation, Hyperparameters tuned (Random)



Gradient Boosting Hiperparameters tuning

```
In [ ]: from sklearn.ensemble import RandomForestClassifier
from sklearn.model_selection import GridSearchCV, RandomizedSearchCV

# Create a Random Forest classifier object
gb_classifier = GradientBoostingClassifier()

# Define the hyperparameter grid
param_grid = {
    'n_estimators': [50, 100, 200],
    'learning_rate': [0.01, 0.1, 0.5],
```

```

    'max_depth': [3, 5, 7],
    'min_samples_split': [2, 5, 10],
    'min_samples_leaf': [1, 2, 4],
    'max_features': ['auto', 'sqrt', 'log2']
}

# Perform randomized search
np.random.seed(42)
random_search = RandomizedSearchCV(estimator=gb_classifier, param_distributions=param_grid, cv=5, n_iter=100)
random_search.fit(features, target)

rand_params = random_search.best_params_

## Perform grid search
#grid_search = GridSearchCV(estimator=gb_classifier, param_grid=param_grid, cv=5, n_iter=100)
#grid_search.fit(features, target)
## Best hyperparameters: {'bootstrap': True, 'max_depth': 20, 'max_features': 'log2', 'min_samples_leaf': 1, 'min_samples_split': 5, 'n_estimators': 100}

## Get the best hyperparameters
#best_params = grid_search.best_params_
#print("Best hyperparameters:", best_params)

best_params = { 'learning_rate': 0.5,
                'max_depth': 7,
                'max_features': 'sqrt',
                'min_samples_leaf': 2,
                'min_samples_split': 10,
                'n_estimators': 50}

best_params_classifier = GradientBoostingClassifier(**best_params)

model_title = "Gradient Boosting Classifier and Cross validation, Hyperparameters tuned"

gb_cv_hp_grid_results, predictions = myCrossValidation(best_params_classifier, skf, features, target)
myPlotROCCurve(target, predictions, model_title)
myPlotConfusionMatrix(target, predictions, model_title)

gb_cv_hp_grid_results.head()

rnd_params_classifier = GradientBoostingClassifier(**rand_params)

model_title = "Gradient Boosting Classifier and Cross validation, Hyperparameters randomized"

gb_cv_hp_rnd_results, predictions = myCrossValidation(rnd_params_classifier, skf, features, target)
myPlotROCCurve(target, predictions, model_title)
myPlotConfusionMatrix(target, predictions, model_title)

gb_cv_hp_rnd_results.head()

```

```
final_results.append(gb_cv_hp_grid_results)
final_results.append(gb_cv_hp_rnd_results)
```

c:\Users\Nemanja\miniconda3\envs\mlenv\lib\site-packages\sklearn\model_selection_validation.py:425: FitFailedWarning:
15 fits failed out of a total of 50.
The score on these train-test partitions for these parameters will be set to nan.
If these failures are not expected, you can try to debug them by setting error_score='raise'.

Below are more details about the failures:

11 fits failed with the following error:

Traceback (most recent call last):

```
File "c:\Users\Nemanja\miniconda3\envs\mlenv\lib\site-packages\sklearn\model_selection\_validation.py", line 732, in _fit_and_score
    estimator.fit(X_train, y_train, **fit_params)
File "c:\Users\Nemanja\miniconda3\envs\mlenv\lib\site-packages\sklearn\base.py", line 1144, in wrapper
    estimator._validate_params()
File "c:\Users\Nemanja\miniconda3\envs\mlenv\lib\site-packages\sklearn\base.py", line 637, in _validate_params
    validate_parameter_constraints(
File "c:\Users\Nemanja\miniconda3\envs\mlenv\lib\site-packages\sklearn\utils\_param_validation.py", line 95, in validate_parameter_constraints
    raise InvalidParameterError(
sklearn.utils._param_validation.InvalidParameterError: The 'max_features' parameter of GradientBoostingClassifier must be an int in the range [1, inf), a float in the range (0.0, 1.0], a str among {'log2', 'sqrt'} or None. Got 'auto' instead.
```

4 fits failed with the following error:


Traceback (most recent call last):

```
File "c:\Users\Nemanja\miniconda3\envs\mlenv\lib\site-packages\sklearn\model_selection\_validation.py", line 732, in _fit_and_score
    estimator.fit(X_train, y_train, **fit_params)
File "c:\Users\Nemanja\miniconda3\envs\mlenv\lib\site-packages\sklearn\base.py", line 1144, in wrapper
    estimator._validate_params()
File "c:\Users\Nemanja\miniconda3\envs\mlenv\lib\site-packages\sklearn\base.py", line 637, in _validate_params
    validate_parameter_constraints(
File "c:\Users\Nemanja\miniconda3\envs\mlenv\lib\site-packages\sklearn\utils\_param_validation.py", line 95, in validate_parameter_constraints
    raise InvalidParameterError(
sklearn.utils._param_validation.InvalidParameterError: The 'max_features' parameter of GradientBoostingClassifier must be an int in the range [1, inf), a float in the range (0.0, 1.0], a str among {'sqrt', 'log2'} or None. Got 'auto' instead.
```


```
warnings.warn(some_fits_failed_message, FitFailedWarning)
```

```
c:\Users\Nemanja\miniconda3\envs\mlenv\lib\site-packages\sklearn\model_selection\_search.py:976: UserWarning: One or more of the test scores are non-finite: [0.97916042
0.95488756 0.98088456 0.97914543 0.98610195          nan
0.9131934          nan 0.95487256          nan]
warnings.warn(
```



Receiver Operating Characteristic (ROC) - Gradient Boosting Classifier and Cross validation, Hyperparameters tuned (Grid)




Gradient Boosting Classifier and Cross validation, Hyperparameters tuned (Grid)



Receiver Operating Characteristic (ROC) - Gradient Boosting Classifier and Cross validation, Hyperparameters tuned (Random)




Gradient Boosting Classifier and Cross validation, Hyperparameters tuned (Random)



Comparison

```
In [ ]: plot_classifier_scores(final_results, feature_names, angle = 45)
```



Zaključujemo da je najbolje rezultate dao Gradient Boosting algoritam sa isprobavanjem svih varijacija parametara, dok je random biranje parametara dalo slične rezultate za oba algoritma.